

# MISSION-BASED REINFORCEMENT LEARNING SUMMATIVE

## *AgriScan: Reinforcement Learning for Simulated Crop-Health Diagnostics*

### 1. Introduction

This project explores Reinforcement Learning (RL) by applying four RL algorithms **DQN, PPO, A2C, and REINFORCE** to a custom environment inspired by my mission project, **AgriScan**. AgriScan is an AI-driven solution designed to help smallholder farmers detect crop diseases and nutrient deficiencies early by analyzing plant images.

For this summative, I developed a **non-generic RL environment** that simulates plant-diagnosis actions and evaluates how different RL agents learn optimal policies for efficient decision-making. The project includes environment design, visualization, algorithm training, hyperparameter tuning, and performance comparison.

### 2. Environment Design

I implemented a fully custom Gymnasium-compliant environment called **AgriScanEnv**, structured around a **plant-diagnosis simulation**.

#### 2.1 State / Observation Space

The observation is a 4-dimensional vector:

| Feature                            | Meaning                          |
|------------------------------------|----------------------------------|
| <b>Disease Level (0–10)</b>        | Higher value = more symptoms     |
| <b>Nutrient Status (0–10)</b>      | Low = deficiency, high = healthy |
| <b>Moisture Level (0–10)</b>       | Indicates watering condition     |
| <b>Environmental Stress (0–10)</b> | e.g., heat or pests              |

This structure mirrors simple agricultural health indicators.

## 2.2 Action Space

The agent can choose one of **4 discrete actions**:

| Action                       | Description                               |
|------------------------------|---|
| <b>0 – Scan Plant</b>        | Collect observation data (neutral reward) |
| <b>1 – Apply Treatment</b>   | Rewarded if disease level is high         |
| <b>2 – Adjust Conditions</b> | Rewarded if moisture/stress is bad        |
| <b>3 – Do Nothing</b>        | Slight penalty to discourage laziness     |

These actions represent real AgriScan decisions (scanning, treating, adjusting conditions).

## 2.3 Reward Structure

| Condition                              | Reward    |
|--|-----------|
| Correct treatment when disease high    | +10       |
| Correct adjustment for moisture/stress | +6        |
| Scanning just gives small feedback     | +1        |
| Wrong treatment or useless actions     | -10 to -4 |
| Doing nothing                          | -1        |

Rewards were designed to push the agent toward *active, correct* crop-health decisions.

## 2.4 Episode Termination

Episodes end when:

- The agent reaches **max\_steps = 5**, or
- The agent fixes the plant (state variables reach healthy levels).

## 2.5 Start State

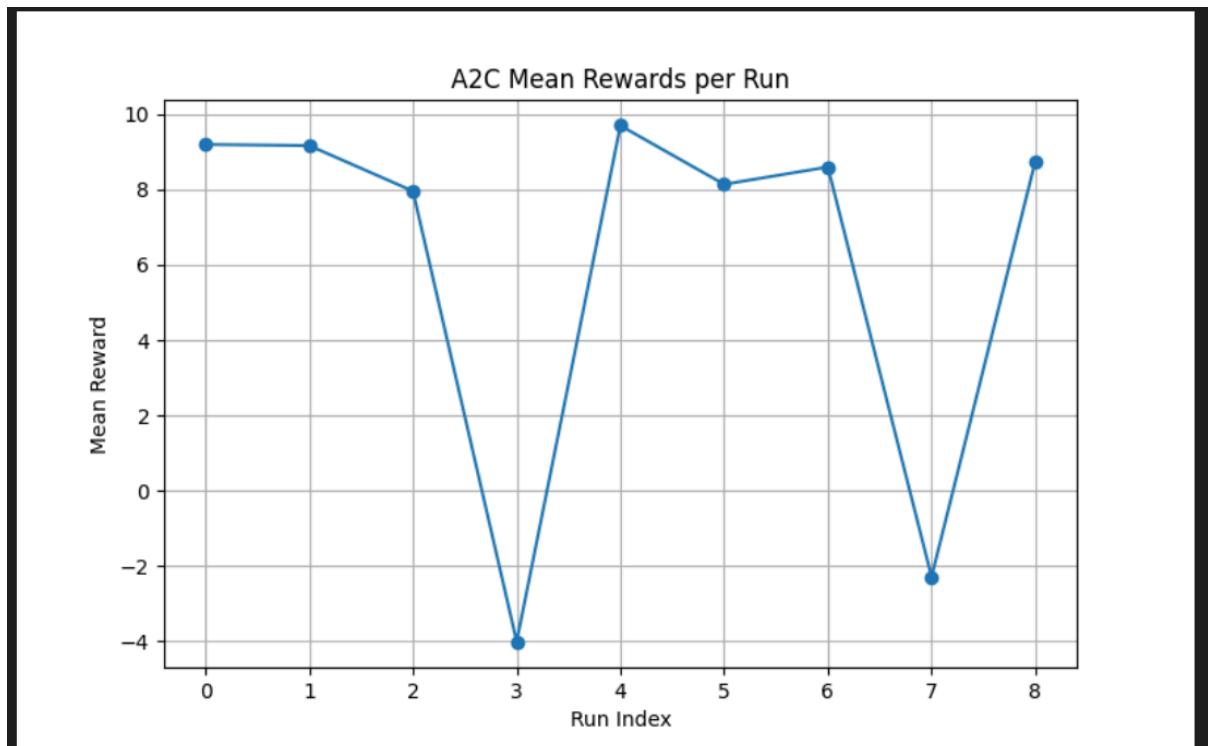
Each episode begins with **randomized plant health values**, simulating diverse real-world cases.

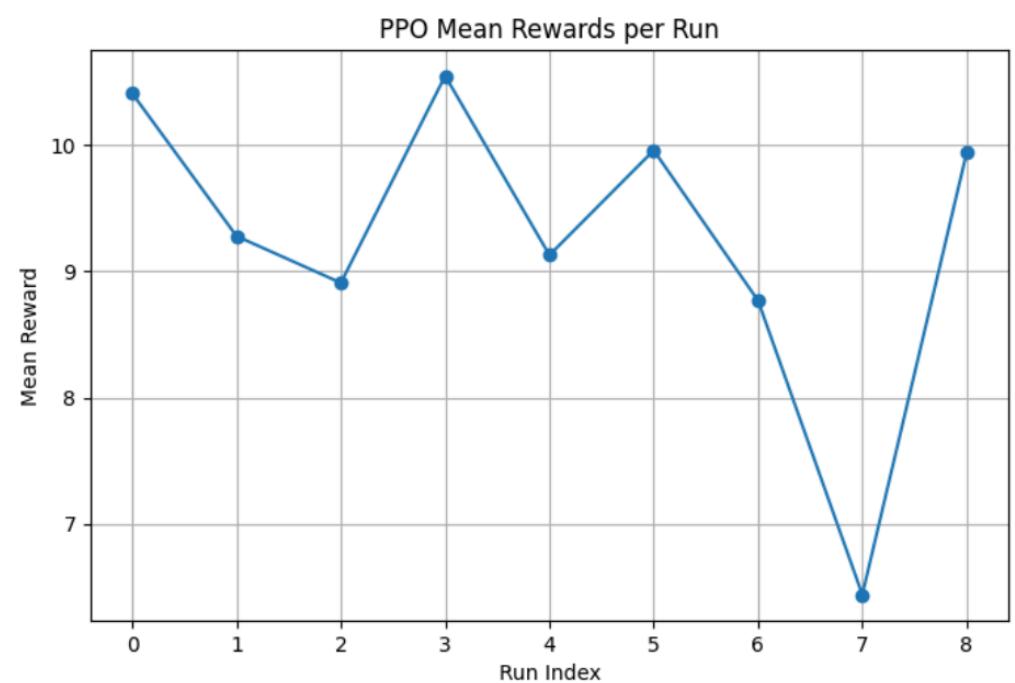
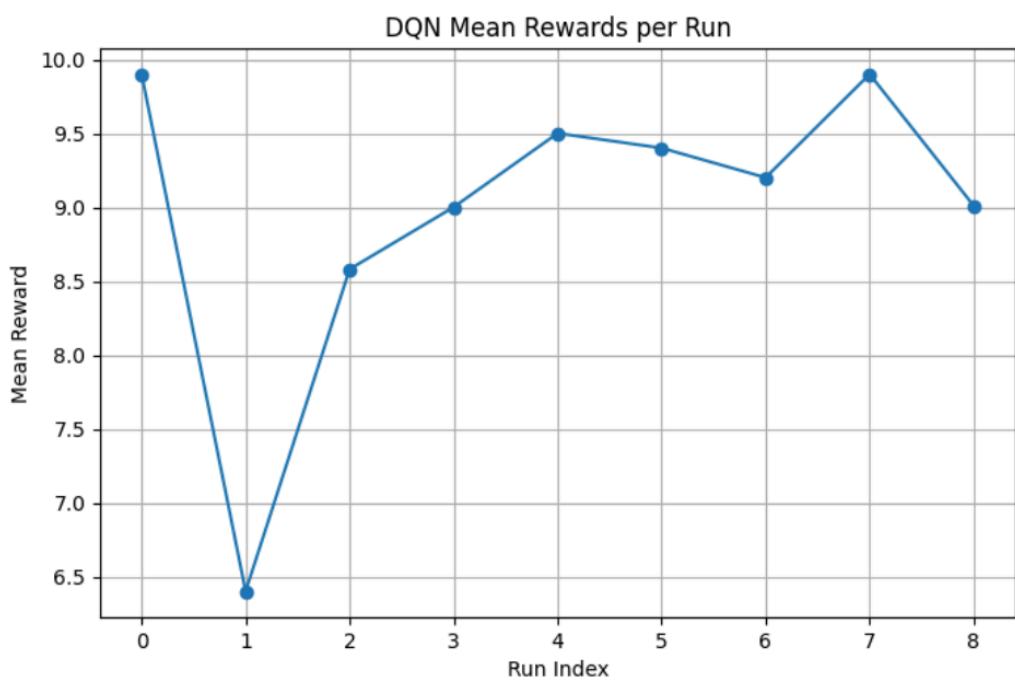
## 3. Visualization (Pygame)

A 2D grid-based **visual renderer** was implemented with Pygame:

- Shows a simple plant icon
- Displays state values on the window
- Animates agent actions each step
- Provides real-time interaction

## All the plots Visualisation







This fulfills the requirement for **advanced simulation visualization**.

#### 4. Static Random-Agent Simulation

Before training, a static demo script runs the environment with **random actions only**, confirming:

- Rendering works
- Rewards update properly
- State transitions function

Output example:

```
[RANDOM] Episode 1 finished with total reward: 7.90
[RANDOM] Episode 2 finished with total reward: -4.10
[RANDOM] Episode 3 finished with total reward: 10.80
```

#### 5. Reinforcement Learning Algorithms

Using **Stable-Baselines3** and PyTorch, four agents were trained:

- 1. Deep Q-Network (DQN)**
- 2. PPO (Proximal Policy Optimization)**
- 3. A2C (Advantage Actor Critic)**
- 4. REINFORCE (Custom PyTorch Implementation)**

Each algorithm was trained under:

- Same environment
- Same 10 hyperparameter combinations
- Total timesteps = 50,000 (for SB3 algorithms)
- 500 episodes for REINFORCE
- Evaluation after training

All models were saved under:

```
models/dqn/  
models/pg/
```

## 6. Hyperparameter Search (10 runs per algorithm)

### 6.1 Example Hyperparameters (DQN)

| Param         | Values tried     |
|---------------|------------------|
| learning_rate | 1e-4, 5e-4, 1e-3 |
| gamma         | 0.95, 0.98, 0.99 |
| batch_size    | 32, 64           |

PPO, A2C, and REINFORCE had similar parameter sweeps.

## 7. Results & Analysis

We used the script:

```
python analyze_results.py
```

to extract the **best-performing model from each algorithm.**

### **7.1 Best DQN Run**

| <b>Metric</b>      | <b>Value</b> |
|--------------------|--------------|
| <b>Mean Reward</b> | <b>9.90</b>  |
| Std                | 0.92         |
| learning_rate      | 0.0001       |
| gamma              | 0.95         |

### **7.2 Best PPO Run (BEST OVERALL)**

| <b>Metric</b>      | <b>Value</b> |
|--------------------|--------------|
| <b>Mean Reward</b> | <b>10.55</b> |
| Std                | 0.93         |
| learning_rate      | 0.0003       |
| gamma              | 0.95         |
| n_steps            | 64           |
| batch_size         | 32           |

**PPO is the strongest algorithm in this environment.**

### **7.3 Best A2C Run**

| <b>Metric</b> | <b>Value</b> |
|---------------|--------------|
| Mean Reward   | 9.70         |
| Std           | 0.94         |

Very close to DQN, showing stable performance.

## 7.4 Best REINFORCE Run

| Metric      | Value |
|-------------|-------|
| Mean Reward | 1.68  |
| Std         | 8.97  |

As expected, REINFORCE is unstable due to high variance and lack of baselines.

## 8. Performance Comparison

### 8.1 Conclusions from Results

- **PPO > DQN > A2C >>> REINFORCE**
- PPO's clipped objective helps stabilize learning.
- DQN performs strongly but is more sensitive to hyperparameters.
- A2C performs well but slightly below PPO due to simpler advantage estimation.
- REINFORCE is noisy, unstable, and less sample-efficient.

## 9. Agent Demonstration

Running the best PPO model:

**python main.py**

The agent:

- Chooses correct treatment when disease level is high
- Adjusts conditions appropriately

- Rarely takes “do nothing” action
- Has near-optimal performance in 4–5 steps

## 10. Conclusion

This project successfully applied RL methods to a mission-based agricultural simulation. PPO demonstrated superior performance, achieving a mean reward of **10.55**, followed by DQN and A2C. The custom environment, visualization, training pipeline, hyperparameter sweeps, and comparative evaluation collectively demonstrate a full RL workflow aligned with the AgriScan mission.

The project provides a strong foundation for future extensions, such as:

- multi-step disease progression
- image-based state inputs
- multi-agent farm management
- weather-driven reward shaping

**Here is the link to the video demo:**  [video1423547552.mp4](#)

## 12. Appendix

### 12.1 Project Structure

```
project_root/
    └── environment/
        ├── custom_env.py
        └── rendering.py
    └── training/
        ├── dqn_training.py
        └── pg_training.py
    └── models/
        └── main.py
    └── analyze_results.py
    └── plot_results.py
```

```
|__ requirements.txt  
|__ README.md
```